

```
# Kelsey Barton
# DSC530-T304
# Final Assignment
# March 2, 2024

# The Dataset: WA_Fn-UseC_-HR-Employee-Attrition

# The Question: Can certain characteristics about a person predict whether
they will choose to remain employed within an organization or leave the
organization?
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load the CSV file into a DataFrame
df = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
# Display the first few rows of the DataFrame
print(df.head())
```

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102		Sales
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	

	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	\
0		1	2 Life Sciences	1		1
1		8	1 Life Sciences	1		2
2		2	2 Other	1		4
3		3	4 Life Sciences	1		5
4		2	1 Medical	1		7

	...	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...		1 80	0	
1	...		4 80	1	
2	...		2 80	0	
3	...		3 80	0	
4	...		4 80	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8		0 1		6
1	10		3 3		10
2	7		3 3		0
3	8		3 3		8

4	6	3	3	2
YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager		
0	4	0	5	
1	7	1	7	
2	0	0	0	
3	7	3	0	
4	2	2	2	

[5 rows x 35 columns]

Describe what the 5 variables mean in the dataset.

This dataset includes many different variables. The most prevalent variables for the experiments question are: age, attrition, rate of pay, travel distance to work, and education. Age represents the age of the employee. Attrition indicates whether the employee has left the company or not. DailyRate, or rate of pay, denotes the daily rate of pay for the employee. DistanceFromHome, or distance to work showcases the distance from the employee's home to the workplace. Education provides the level of education attained by the employee. These are the five main variables being used in the provided code for analysis and visualization.

Include a histogram of each of the 5 variables.

Histograms

plt.figure(figsize=(15, 10))

Age

plt.subplot(2, 3, 1)
plt.hist(df['Age'], bins=20, edgecolor='black')
plt.title('Age')

Attrition

plt.subplot(2, 3, 2)
plt.hist(df['Attrition'], bins=2, edgecolor='black')
plt.title('Attrition')

DailyRate

plt.subplot(2, 3, 3)
plt.hist(df['DailyRate'], bins=20, edgecolor='black')
plt.title('DailyRate')

DistanceFromHome

plt.subplot(2, 3, 4)
plt.hist(df['DistanceFromHome'], bins=20, edgecolor='black')
plt.title('DistanceFromHome')

Education

plt.subplot(2, 3, 5)

```
plt.hist(df['Education'], bins=5, edgecolor='black')
plt.title('Education')

plt.tight_layout()
plt.show()

plt.figure(figsize=(12, 10))

plt.subplot(3, 3, 1)
plt.hist(df['Age'], bins=20, edgecolor='black')
plt.title('Age')

plt.subplot(3, 3, 2)
attrition_counts = df['Attrition'].value_counts()
plt.bar(attrition_counts.index, attrition_counts.values, color='skyblue')
plt.title('Attrition')
plt.xticks(rotation=45)

plt.subplot(3, 3, 3)
plt.hist(df['DailyRate'], bins=20, edgecolor='black')
plt.title('DailyRate')

plt.subplot(3, 3, 4)
plt.hist(df['DistanceFromHome'], bins=20, edgecolor='black')
plt.title('DistanceFromHome')

plt.subplot(3, 3, 5)
education_counts = df['Education'].value_counts()
plt.bar(education_counts.index, education_counts.values, color='skyblue')
plt.title('Education')
plt.xticks(rotation=45)

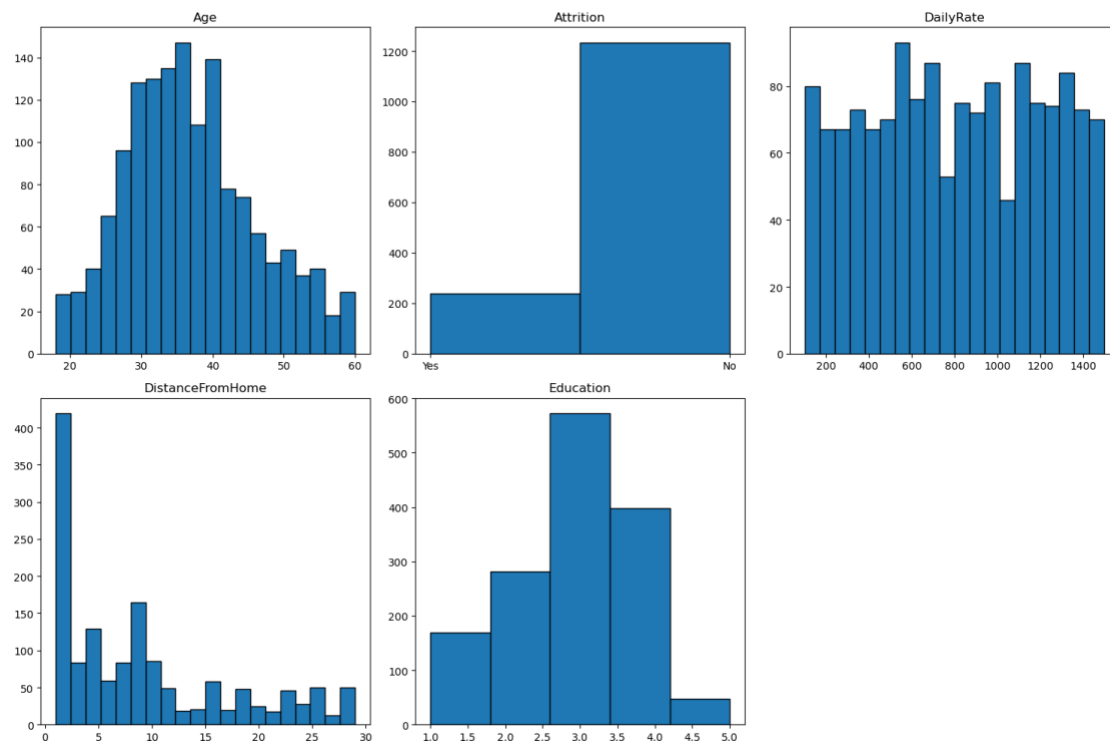
plt.subplot(3, 3, 6)
plt.hist(df['MonthlyIncome'], bins=20, edgecolor='black')
plt.title('MonthlyIncome')

plt.subplot(3, 3, 7)
plt.hist(df['PercentSalaryHike'], bins=20, edgecolor='black')
plt.title('PercentSalaryHike')

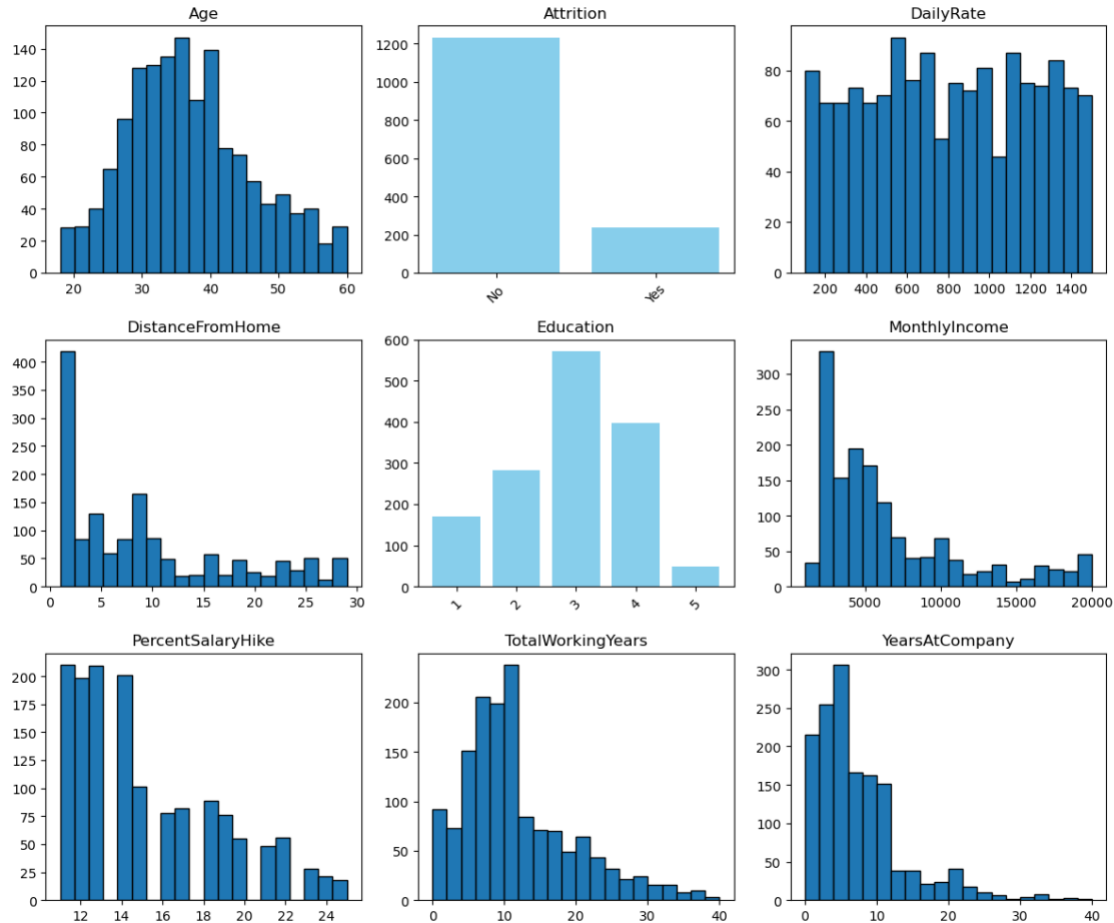
plt.subplot(3, 3, 8)
plt.hist(df['TotalWorkingYears'], bins=20, edgecolor='black')
plt.title('TotalWorkingYears')

plt.subplot(3, 3, 9)
plt.hist(df['YearsAtCompany'], bins=20, edgecolor='black')
plt.title('YearsAtCompany')
```

```
plt.tight_layout()
plt.show()
```



png



png

Identify any outliers and explain the reasoning for them being outliers and how you believe they should be handled.

The outliers for age are any data points that fall far from the bulk of the data, such as ages below 20 or above 60. Outliers in age might represent valid data points (e.g., experienced employees) or data entry errors.

The outliers for daily rate are any extremely high or low daily rates that deviate significantly from the majority of the data. Similar to age, outliers in daily rate could be due to genuine differences in salaries or data entry errors. If they are legitimate data points, they should be kept in the analysis. Otherwise, they might be removed or adjusted after investigating the reasons behind them.

The distance from home outliers are large distances from home that are not representative of the majority of employees. These outliers could indicate remote employees, employees who have relocated, or data entry errors. If they are valid data points, they should be retained; otherwise, they might be adjusted or removed.

The education outliers consist of unusual education levels compared to the majority of employees. These outliers might indicate employees with advanced degrees or employees with less formal education than the majority. They might

be retained, grouped differently, or treated separately depending on how the relate to the question.

Include the other descriptive characteristics about the variables: mean, mode, spread, and tails.

Descriptive statistics

```
print("Descriptive Statistics:")
```

```
print(df.describe())
```

```
descriptive_stats = {
    'Variable': ['Age', 'Attrition', 'DailyRate', 'DistanceFromHome',
    'Education'],
    'Mean': [df['Age'].mean(), None, df['DailyRate'].mean(),
    df['DistanceFromHome'].mean(), df['Education'].mean()],
    'Mode': [df['Age'].mode().values[0], df['Attrition'].mode().values[0],
    df['DailyRate'].mode().values[0], df['DistanceFromHome'].mode().values[0],
    df['Education'].mode().values[0]],
    'Spread (Standard Deviation)': [df['Age'].std(), None,
    df['DailyRate'].std(), df['DistanceFromHome'].std(), df['Education'].std()],
    'Tails': ['Two-tailed', None, 'Right-tailed', 'Right-tailed', None]
}
```

Convert the dictionary to a DataFrame

```
descriptive_df = pd.DataFrame(descriptive_stats)
```

Print descriptive statistics

```
print(descriptive_df)
```

Descriptive Statistics:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount
\					
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0
mean	36.923810	802.485714	9.192517	2.912925	1.0
std	9.135373	403.509100	8.106864	1.024165	0.0
min	18.000000	102.000000	1.000000	1.000000	1.0
25%	30.000000	465.000000	2.000000	2.000000	1.0
50%	36.000000	802.000000	7.000000	3.000000	1.0
75%	43.000000	1157.000000	14.000000	4.000000	1.0
max	60.000000	1499.000000	29.000000	5.000000	1.0

	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement
\				
count	1470.000000	1470.000000	1470.000000	1470.000000
mean	1024.865306	2.721769	65.891156	2.729932
std	602.024335	1.093082	20.329428	0.711561
min	1.000000	1.000000	30.000000	1.000000
25%	491.250000	2.000000	48.000000	2.000000
50%	1020.500000	3.000000	66.000000	3.000000
75%	1555.750000	4.000000	83.750000	3.000000

max	2068.000000	4.000000	100.000000	4.000000
-----	-------------	----------	------------	----------

	JobLevel	...	RelationshipSatisfaction	StandardHours	\
count	1470.000000	...	1470.000000	1470.0	
mean	2.063946	...	2.712245	80.0	
std	1.106940	...	1.081209	0.0	
min	1.000000	...	1.000000	80.0	
25%	1.000000	...	2.000000	80.0	
50%	2.000000	...	3.000000	80.0	
75%	3.000000	...	4.000000	80.0	
max	5.000000	...	4.000000	80.0	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
count	1470.000000	1470.000000	1470.000000	
mean	0.793878	11.279592	2.799320	
std	0.852077	7.780782	1.289271	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
count	1470.000000	1470.000000	1470.000000	
mean	2.761224	7.008163	4.229252	
std	0.706476	6.126525	3.623137	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	9.000000	7.000000	
max	4.000000	40.000000	18.000000	

	YearsSinceLastPromotion	YearsWithCurrManager
count	1470.000000	1470.000000
mean	2.187755	4.123129
std	3.222430	3.568136
min	0.000000	0.000000
25%	0.000000	2.000000
50%	1.000000	3.000000
75%	3.000000	7.000000
max	15.000000	17.000000

[8 rows x 26 columns]

	Variable	Mean	Mode	Spread (Standard Deviation)	\
0	Age	36.923810	35	9.135373	
1	Attrition	NaN	No	NaN	
2	DailyRate	802.485714	691	403.509100	
3	DistanceFromHome	9.192517	2	8.106864	
4	Education	2.912925	3	1.024165	

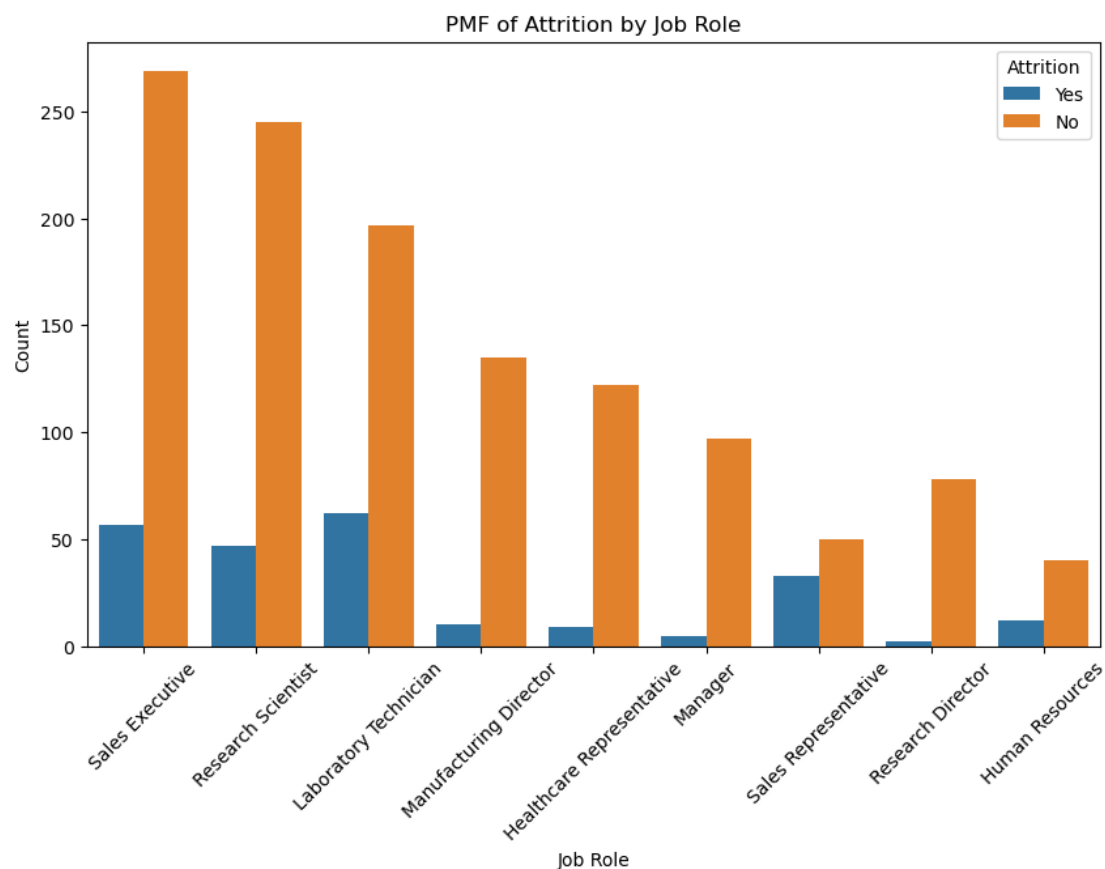
```

Tails
0    Two-tailed
1      None
2  Right-tailed
3  Right-tailed
4      None

# Compare two scenarios in your data using a PMF.

# PMF comparing attrition for different job roles
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='JobRole', hue='Attrition')
plt.title('PMF of Attrition by Job Role')
plt.xticks(rotation=45)
plt.xlabel('Job Role')
plt.ylabel('Count')
plt.legend(title='Attrition')
plt.show()

```



png

```


# Create 1 CDF with one of your variables.

```



```
# CDF of MonthlyIncome
sorted_income = df['MonthlyIncome'].sort_values()
cdf = np.arange(1, len(sorted_income) + 1) / len(sorted_income)

plt.figure(figsize=(8, 6))
plt.plot(sorted_income, cdf, marker='.', linestyle='none')
plt.xlabel('Monthly Income')
plt.ylabel('CDF')
plt.title('CDF of Monthly Income')
plt.grid(True)
plt.show()
```



png

What does this tell you about your variable and how does it address the question you are trying to answer?

The CTF of monthly income provides insight into the distribution of monthly income within the organization. From left to right on the graph above, the CDF curve rises, indicating an increase in the cumulative probability. The corresponding Y value represents the proportion of individuals in the data set with a monthly income less than or equal to that value. The graph is providing information on the distribution of monthly income. The graph is showing that less than 20% of the organization is making less than \$2500 a

month. It also shows that about half of the organization is making \$5000 a month. There is a significant increase in the second half of the organization. This increase is anywhere from over \$5000 a month to over \$20,000 a month. It appears that the highest paid individuals within the organization make anywhere from \$10,000 a month to \$20,000 a month. The top 10% within the organization have a \$5000 differential of \$15,000-\$20,000. This specific variable addresses the question of likelihood to stay within an organization because people are driven by money. This is not always the circumstance, but around 70% of this organization is making less than \$6000 a month. That is anywhere from \$500-\$1500 per week. The most highly paid person within the organization is making \$20,000. That is a \$14,000 differential at \$5000 a week. If the employees within this organization know about this differential and do not have a way of earning higher compensation, that could be an incentive to leave or seek another job.

Plot 1 analytical distribution.

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import norm
from scipy import stats
```

Load the CSV file into a DataFrame

```
df = pd.read_csv("WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

Select the 'Age' variable for analysis

```
data = df['Age']
```

Fit a normal distribution to the data

```
mu, std = norm.fit(data)
```

Plot the histogram of the data

```
plt.figure(figsize=(8, 6))
sns.histplot(data, kde=True, color='skyblue', stat='density', bins=20)
```


Plot the fitted normal distribution

```
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)
plt.plot(x, p, 'k', linewidth=2)
```

Add labels and title

```
plt.xlabel('Age')
plt.ylabel('Density')
plt.title('Histogram with Fitted Normal Distribution (Age)')
plt.legend(['Fitted Normal Distribution', 'Empirical Data'])
```

```
plt.show()
```



png

Provide your analysis on how it applies to the dataset you have chosen.

The age histogram shown above showcases the age range and density of age within the organization. Within the histogram, it can be seen that the ages range from 18 years old to 60 years old throughout the organization. The most popular age for someone working within the organization is around 36 years of age. The empirical data line shows that the median age is right around that 36 mark as well. The fitted normal distribution is skewed to the left showing that most of the employees within the organization are 40 years or below. It is much less common to see employees above the age of 40 and none at all above the age of 60.

Create two scatter plots comparing two variables and provide your analysis on correlation and causation.

```
import seaborn as sns
```

```
# Selecting the variables
```

```
x = df['Age']
```

```
y = df['MonthlyIncome']
```

```
# Scatter plot
```

```
plt.figure(figsize=(8, 6))
```

```

sns.scatterplot(x=x, y=y, color='skyblue')
plt.title('Scatter Plot of Age vs. Monthly Income')
plt.xlabel('Age')
plt.ylabel('Monthly Income')
plt.grid(True)
plt.show()


# Covariance
covariance = np.cov(x, y)[0][1]
print("Covariance:", covariance)

# Pearson's correlation coefficient
correlation = np.corrcoef(x, y)[0][1]
print("Pearson's correlation coefficient:", correlation)

# Non-Linear Relationships

# Jointplot with regression line
sns.jointplot(x=x, y=y, kind="reg", color='skyblue')
plt.xlabel('Age')
plt.ylabel('Monthly Income')
plt.show()


```



png

Covariance: 21412.198982138805

Pearson's correlation coefficient: 0.4978545669265806



png

Conduct a test on your hypothesis.

```
from scipy.stats import ttest_ind
```

Split the data into two groups based on attrition

```
attrition_yes = df[df['Attrition'] == 'Yes']['MonthlyIncome']
```

```
attrition_no = df[df['Attrition'] == 'No']['MonthlyIncome']
```

Perform the t-test

```
t_statistic, p_value = ttest_ind(attrition_yes, attrition_no)
```

Print the results

```
print("T-Statistic:", t_statistic)
```

```

print("P-Value:", p_value)

# Interpretation
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis. There is a significant difference in
monthly income between employees who have left the company and those who are
still with the company.")
else:
    print("Fail to reject the null hypothesis. There is no significant
difference in monthly income between employees who have left the company and
those who are still with the company.")

T-Statistic: -6.203935765608938
P-Value: 7.14736398535381e-10
Reject the null hypothesis. There is a significant difference in monthly
income between employees who have left the company and those who are still
with the company.

# Conduct a regression analysis on either one dependent and one explanatory
variable, or multiple explanatory variables.

import statsmodels.api as sm

# Add a constant term to the explanatory variable
X = sm.add_constant(df['Age'])
y = df['MonthlyIncome']

# Fit the linear regression model
model = sm.OLS(y, X).fit()

# Print the regression summary
print(model.summary())

```

```

=====
                        OLS Regression Results
=====
=
Dep. Variable:          MonthlyIncome    R-squared:
0.248
Model:                  OLS              Adj. R-squared:
0.247
Method:                 Least Squares    F-statistic:
483.8
Date:                   Mon, 26 Feb 2024  Prob (F-statistic):      6.67e-
93
Time:                   23:52:07         Log-Likelihood:          -
14308.
No. Observations:       1470             AIC:
2.862e+04
Df Residuals:           1468             BIC:

```

```

2.863e+04
Df Model: 1
Covariance Type: nonrobust
=====
=
coef std err t P>|t| [0.025
0.975]
-----
-
const -2970.6712 443.702 -6.695 0.000 -3841.030 -
2100.313
Age 256.5716 11.665 21.995 0.000 233.689
279.454
=====
=
Omnibus: 140.178 Durbin-Watson:
2.069
Prob(Omnibus): 0.000 Jarque-Bera (JB):
182.119
Skew: 0.799 Prob(JB): 2.84e-
40
Kurtosis: 3.649 Cond. No.
159.
=====
=

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.

```